

## Task-D: Collinear features and their effect on linear models

```
In [6]: %matplotlib inline
import warnings
warnings.filterwarnings("ignore")
import pandas as pd
import numpy as np
from sklearn.datasets import load_iris
from sklearn.linear_model import SGDClassifier
from sklearn.model_selection import GridSearchCV
import seaborn as sns
import matplotlib.pyplot as plt
```

```
In [7]: data = pd.read_csv("C:\\Users\\Akashraj D S\\Downloads\\8_LinearModels-20210710T")
```

```
In [8]: data.head()
```

```
Out[8]:
```

	x	y	z	x*x	2*y	2*z+3*x*x	w	target
0	-0.581066	0.841837	-1.012978	-0.604025	0.841837	-0.665927	-0.536277	0
1	-0.894309	-0.207835	-1.012978	-0.883052	-0.207835	-0.917054	-0.522364	0
2	-1.207552	0.212034	-1.082312	-1.150918	0.212034	-1.166507	0.205738	0
3	-1.364174	0.002099	-0.943643	-1.280666	0.002099	-1.266540	-0.665720	0
4	-0.737687	1.051772	-1.012978	-0.744934	1.051772	-0.792746	-0.735054	0

```
In [9]: X = data.drop(['target'], axis=1).values
Y = data['target'].values
```

### Doing perturbation test to check the presence of collinearity

#### Task: 1 Logistic Regression

1. **Finding the Correlation between the features**
  - a. check the correlation between the features
  - b. plot heat map of correlation matrix using seaborn heatmap
2. **Finding the best model for the given data**
  - a. Train Logistic regression on data(X,Y) that we have created in the above cell
  - b. Find the best hyper parameter alpha with hyper parameter tuning using k-fold cross validation (grid search CV or random search CV make sure you choose the alpha in log space)
  - c. Create a new Logistic regression with the best alpha (search for how to get the best hyper parameter value), name the best model as 'best\_model'
3. **Getting the weights with the original data**
  - a. train the 'best\_model' with X, Y
  - b. Check the accuracy of the model 'best\_model\_accuracy'
  - c. Get the weights W using best\_model.coef\_
4. **Modifying original data**
  - a. Add a noise(order of  $10^{-2}$ ) to each element of X and get the new data set X' ( $X' = X + e$ )
  - b. Train the same 'best\_model' with data (X', Y)
  - c. Check the accuracy of the model 'best\_model\_accuracy\_edited'
  - d. Get the weights W' using best\_model.coef\_
5. **Checking deviations in metric and weights**
  - a. find the difference between 'best\_model\_accuracy\_edited' and 'best\_model\_accuracy'
  - b. find the absolute change between each value of W and W' ==>  $|W - W'|$
  - c. print the top 4 features which have higher % change in weights compare to the other feature

## Task: 2 Linear SVM

1. Do the same steps (2, 3, 4, 5) we have done in the above task 1.

**Do write the observations based on the results you get from the deviations of weights in both Logistic Regression and linear SVM**

## Finding the Correlation between the features

```
In [10]: correlation = data.corr()
```


```
In [11]: correlation
```

```
Out[11]:
```

	x	y	z	x*x	2*y	2*z+3*x*x	w	target
x	1.000000	-0.205926	0.812458	0.997947	-0.205926	0.996252	0.583277	0.728290
y	-0.205926	1.000000	-0.602663	-0.209289	1.000000	-0.261123	-0.401790	-0.690684
z	0.812458	-0.602663	1.000000	0.807137	-0.602663	0.847163	0.674486	0.969990
x*x	0.997947	-0.209289	0.807137	1.000000	-0.209289	0.997457	0.583803	0.719570
2*y	-0.205926	1.000000	-0.602663	-0.209289	1.000000	-0.261123	-0.401790	-0.690684
2*z+3*x*x	0.996252	-0.261123	0.847163	0.997457	-0.261123	1.000000	0.606860	0.764729
w	0.583277	-0.401790	0.674486	0.583803	-0.401790	0.606860	1.000000	0.641750
target	0.728290	-0.690684	0.969990	0.719570	-0.690684	0.764729	0.641750	1.000000

```
In [12]: sns.heatmap(correlation,annot=True)
```

```
Out[12]: <matplotlib.axes._subplots.AxesSubplot at 0x28b4d7b0408>
```



```
In [13]: #Y and 2Y are highly correlated
#X and X*X are highly correlated
#2*Z+3*X*X is highly correlated with X and X*X
```

## Logistic Regression

### Finding the best model for given data

```
In [14]: import numpy
alpha = numpy.logspace(0.001,1,5)
```

```
In [15]: alpha
```

```
Out[15]: array([ 1.00230524,  1.78135304,  3.16592046,  5.62665128, 10.          ])
```

```
In [16]: params={'alpha':alpha}
clf=SGDClassifier(loss='log')
cross_val = GridSearchCV(estimator=clf,param_grid = params, cv=5)
cross_val.fit(X,Y)
```

```
Out[16]: GridSearchCV(cv=5, error_score='raise-deprecating',
                      estimator=SGDClassifier(alpha=0.0001, average=False,
                                              class_weight=None, early_stopping=False,
                                              epsilon=0.1, eta0=0.0, fit_intercept=True,
                                              l1_ratio=0.15, learning_rate='optimal',
                                              loss='log', max_iter=1000,
                                              n_iter_no_change=5, n_jobs=None,
                                              penalty='l2', power_t=0.5,
                                              random_state=None, shuffle=True, tol=0.001,
                                              validation_fraction=0.1, verbose=0,
                                              warm_start=False),
                      iid='warn', n_jobs=None,
                      param_grid={'alpha': array([ 1.00230524,  1.78135304,  3.16592046,
6,  5.62665128, 10.          ])}),
                      pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
                      scoring=None, verbose=0)
```

```
In [17]: pd.DataFrame(cross_val.cv_results_)
```

```
Out[17]:
```

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_alpha	param
0	0.002401	0.001854	0.000456	0.000456	1.00231	{'alpha': 1.0023052380778996}
1	0.001001	0.000002	0.000256	0.000392	1.78135	{'alpha': 1.7813530430086404}
2	0.000601	0.000491	0.000401	0.000491	3.16592	{'alpha': 3.1659204634322378}
3	0.000601	0.000491	0.000400	0.000490	5.62665	{'alpha': 5.626651280675067}
4	0.000738	0.000388	0.000198	0.000397	10	{'alpha': 10.000000000000001}

```
In [18]: best_model_lr = cross_val.best_estimator_
```

## Getting the weights with original data

```
In [19]: best_model_lr.fit(X,Y)
```

```
Out[19]: SGDClassifier(alpha=1.0023052380778996, average=False, class_weight=None,
    early_stopping=False, epsilon=0.1, eta0=0.0, fit_intercept=True,
    l1_ratio=0.15, learning_rate='optimal', loss='log', max_iter=100
    0,
    n_iter_no_change=5, n_jobs=None, penalty='l2', power_t=0.5,
    random_state=None, shuffle=True, tol=0.001,
    validation_fraction=0.1, verbose=0, warm_start=False)
```

```
In [20]: y_pred_lr = best_model_lr.predict(X)
```

```
In [21]: from sklearn.metrics import accuracy_score
best_model_accuracy = accuracy_score(Y,y_pred_lr)
best_model_accuracy
```

```
Out[21]: 0.99
```

```
In [22]: weights_without_noise = best_model_lr.coef_
```

## Modifying original data

```
In [23]: X_dash = data.drop(['target'], axis=1)
X_dash = X_dash.add(0.01)
X_dash = X_dash.values
```

```
In [24]: best_model_lr.fit(X_dash,Y)
```

```
Out[24]: SGDClassifier(alpha=1.0023052380778996, average=False, class_weight=None,
    early_stopping=False, epsilon=0.1, eta0=0.0, fit_intercept=True,
    l1_ratio=0.15, learning_rate='optimal', loss='log', max_iter=100
    0,
    n_iter_no_change=5, n_jobs=None, penalty='l2', power_t=0.5,
    random_state=None, shuffle=True, tol=0.001,
    validation_fraction=0.1, verbose=0, warm_start=False)
```

```
In [25]: y_pred_lr_with_noise = best_model_lr.predict(X_dash)
```

```
In [26]: best_model_accuracy_edited = accuracy_score(Y,y_pred_lr_with_noise)
best_model_accuracy_edited
```

```
Out[26]: 1.0
```

```
In [27]: weights_with_noise = best_model_lr.coef_
```

## Checking deviations in metric and weights

```
In [28]: difference_accuracy = best_model_accuracy - best_model_accuracy_edited
difference_accuracy
```

```
Out[28]: -0.010000000000000009
```

```
In [29]: values = abs(weights_without_noise - weights_with_noise)
```

```
In [30]: values
```

```
Out[30]: array([[0.00385742, 0.00116328, 0.00032523, 0.00447727, 0.00116328,
                0.00407881, 0.00054663]])
```

```
In [31]: keys = data.columns[0:len(data.columns)-1]
```

```
In [32]: keys
```

```
Out[32]: Index(['x', 'y', 'z', 'x*x', '2*y', '2*z+3*x*x', 'w'], dtype='object')
```

```
In [33]: d={}
for i in range(0,len(values[0])):
    d[keys[i]]=values[0][i]
d
```

```
Out[33]: {'x': 0.0038574172414173324,
'y': 0.0011632779234296842,
'z': 0.0003252314872647655,
'x*x': 0.004477268693702935,
'2*y': 0.0011632779234296842,
'2*z+3*x*x': 0.004078808976586579,
'w': 0.0005466275841054802}
```

```
In [34]: sorted_by_feature = sorted(d.items(), key = lambda kv: kv[1], reverse=True)
```

```
In [35]: sorted_by_feature
```

```
Out[35]: [('x*x', 0.004477268693702935),
('2*z+3*x*x', 0.004078808976586579),
('x', 0.0038574172414173324),
('y', 0.0011632779234296842),
('2*y', 0.0011632779234296842),
('w', 0.0005466275841054802),
('z', 0.0003252314872647655)]
```

Type *Markdown* and LaTeX:  $\alpha^2$

## SVM

### Finding the best model for given data

```
In [36]: import numpy
alpha_svm = numpy.logspace(0.001,1,5)
```

```
In [37]: alpha
```

```
Out[37]: array([ 1.00230524,  1.78135304,  3.16592046,  5.62665128, 10.          ])
```

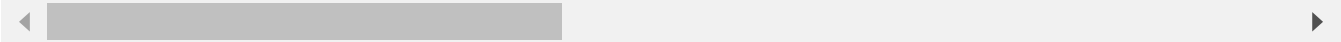
```
In [38]: params={'alpha':alpha}
clf=SGDClassifier(loss='hinge')
cross_val_svm = GridSearchCV(estimator=clf,param_grid = params, cv=5)
cross_val_svm.fit(X,Y)
```

```
Out[38]: GridSearchCV(cv=5, error_score='raise-deprecating',
                      estimator=SGDClassifier(alpha=0.0001, average=False,
                                              class_weight=None, early_stopping=False,
                                              epsilon=0.1, eta0=0.0, fit_intercept=True,
                                              l1_ratio=0.15, learning_rate='optimal',
                                              loss='hinge', max_iter=1000,
                                              n_iter_no_change=5, n_jobs=None,
                                              penalty='l2', power_t=0.5,
                                              random_state=None, shuffle=True, tol=0.0001,
                                              validation_fraction=0.1, verbose=0,
                                              warm_start=False),
                      iid='warn', n_jobs=None,
                      param_grid={'alpha': array([ 1.00230524,  1.78135304,  3.16592046,
5.62665128, 10.          ])},
                      pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
                      scoring=None, verbose=0)
```

```
In [39]: pd.DataFrame(cross_val_svm.cv_results_)
```

```
Out[39]:
```

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_alpha	param
0	0.001690	0.000588	0.000938	0.000539	1.00231	{'alpha': 1.0023052380778996}
1	0.001191	0.000394	0.000402	0.000493	1.78135	{'alpha': 1.7813530430086404}
2	0.001648	0.000447	0.000056	0.000113	3.16592	{'alpha': 3.1659204634322378}
3	0.001040	0.000077	0.000000	0.000000	5.62665	{'alpha': 5.626651280675067}
4	0.000600	0.000490	0.000210	0.000420	10	{'alpha': 10.0}



```
In [40]: best_model_svm = cross_val_svm.best_estimator_
```

## Getting the weights with original data

```
In [41]: best_model_svm.fit(X,Y)
```

```
Out[41]: SGDClassifier(alpha=1.0023052380778996, average=False, class_weight=None,
    early_stopping=False, epsilon=0.1, eta0=0.0, fit_intercept=True,
    l1_ratio=0.15, learning_rate='optimal', loss='hinge',
    max_iter=1000, n_iter_no_change=5, n_jobs=None, penalty='l2',
    power_t=0.5, random_state=None, shuffle=True, tol=0.001,
    validation_fraction=0.1, verbose=0, warm_start=False)
```

```
In [42]: y_pred_svm = best_model_svm.predict(X)
```

```
In [43]: from sklearn.metrics import accuracy_score
best_model_accuracy_svm = accuracy_score(Y,y_pred_svm)
best_model_accuracy_svm
```

```
Out[43]: 1.0
```

```
In [44]: weights_without_noise_svm = best_model_svm.coef_
```

## Modifying original data

```
In [45]: X_dash = data.drop(['target'], axis=1)
X_dash = X_dash.add(0.01)
X_dash_svm = X_dash.values
```

```
In [46]: best_model_svm.fit(X_dash_svm,Y)
```

```
Out[46]: SGDClassifier(alpha=1.0023052380778996, average=False, class_weight=None,
    early_stopping=False, epsilon=0.1, eta0=0.0, fit_intercept=True,
    l1_ratio=0.15, learning_rate='optimal', loss='hinge',
    max_iter=1000, n_iter_no_change=5, n_jobs=None, penalty='l2',
    power_t=0.5, random_state=None, shuffle=True, tol=0.001,
    validation_fraction=0.1, verbose=0, warm_start=False)
```

```
In [47]: y_pred_svm_with_noise = best_model_svm.predict(X_dash_svm)
```

```
In [48]: best_model_accuracy_svm_edited = accuracy_score(Y,y_pred_svm_with_noise)
best_model_accuracy_svm_edited
```

```
Out[48]: 1.0
```

```
In [49]: weights_with_noise_svm = best_model_svm.coef_
```

## Checking deviations in metric and weights

```
In [50]: difference_accuracy_svm = best_model_accuracy_svm - best_model_accuracy_svm_edit
difference_accuracy_svm
```

```
Out[50]: 0.0
```



```
In [51]: values_svm = abs(weights_without_noise_svm - weights_with_noise_svm)
```

```
In [52]: d_svm={}
for i in range(0,len(values_svm[0])):
    d_svm[keys[i]]=values_svm[0][i]
d_svm
```

```
Out[52]: {'x': 0.00085927662336438,
'y': 0.0010573313722122002,
'z': 0.00019098094162989243,
'x*x': 0.0013565964789609142,
'2*y': 0.0010573313722122002,
'2*z+3*x*x': 0.0012425583672647234,
'w': 0.001882516541764967}
```

```
In [53]: sorted_by_feature_svm = sorted(d_svm.items(), key = lambda kv: kv[1], reverse=True)
sorted_by_feature_svm
```

```
Out[53]: [('w', 0.001882516541764967),
('x*x', 0.0013565964789609142),
('2*z+3*x*x', 0.0012425583672647234),
('y', 0.0010573313722122002),
('2*y', 0.0010573313722122002),
('x', 0.00085927662336438),
('z', 0.00019098094162989243)]
```

## Observations

If the features are said to be non-collinear, there should not be any differences between them. In both LR and SVM, it shows difference between weights. So they are collinear

```
In [ ]:
```