

Implement SGD Classifier with Logloss and L2 regularization Using SGD without using sklearn

There will be some functions that start with the word "grader" ex: grader_weights(), grader_sigmoid(), grader_logloss() etc, you should not change those function definition.

Every Grader function has to return True.

Importing packages

```
In [1]: import numpy as np
import pandas as pd
from sklearn.datasets import make_classification
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn import linear_model
```

Creating custom dataset

```
In [2]: # please don't change random_state
X, y = make_classification(n_samples=50000, n_features=15, n_informative=10,
n_redundant=5,
                                n_classes=2, weights=[0.7], class_sep=0.7, random_
state=15)
# make_classification is used to create custom dataset
# Please check this link (https://scikit-learn.org/stable/modules/generated/sklearn.datasets.make\_classification.html) for more details
```

```
In [3]: X.shape, y.shape
```

```
Out[3]: ((50000, 15), (50000,))
```

Splitting data into train and test

```
In [4]: #please don't change random state
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, ran
dom_state=15)
```

```
In [5]: # Standardizing the data.
scaler = StandardScaler()
x_train = scaler.fit_transform(X_train)
x_test = scaler.transform(X_test)
```

```
In [6]: X_train.shape, y_train.shape, X_test.shape, y_test.shape
```

```
Out[6]: ((37500, 15), (37500,)), (12500, 15), (12500,))
```

SGD classifier

```
In [7]: # alpha : float
# Constant that multiplies the regularization term.

# eta0 : double
# The initial learning rate for the 'constant', 'invscaling' or 'adaptive' schedules.

clf = linear_model.SGDClassifier(eta0=0.0001, alpha=0.0001, loss='log', random_state=15, penalty='l2', tol=1e-3, verbose=2, learning_rate='constant')
clf
# Please check this documentation (https://scikit-learn.org/stable/modules/generated/sklearn.linear\_model.SGDClassifier.html)
```

```
Out[7]: SGDClassifier(alpha=0.0001, average=False, class_weight=None,
    early_stopping=False, epsilon=0.1, eta0=0.0001,
    fit_intercept=True, l1_ratio=0.15, learning_rate='constant',
    loss='log', max_iter=1000, n_iter_no_change=5, n_jobs=None,
    penalty='l2', power_t=0.5, random_state=15, shuffle=True,
    tol=0.001, validation_fraction=0.1, verbose=2, warm_start=False)
e)
```

```
In [8]: clf.fit(X=X_train, y=y_train) # fitting our model
```

```
-- Epoch 1
Norm: 0.77, NNZs: 15, Bias: -0.316653, T: 37500, Avg. loss: 0.455552
Total training time: 0.02 seconds.
-- Epoch 2
Norm: 0.91, NNZs: 15, Bias: -0.472747, T: 75000, Avg. loss: 0.394686
Total training time: 0.03 seconds.
-- Epoch 3
Norm: 0.98, NNZs: 15, Bias: -0.580082, T: 112500, Avg. loss: 0.385711
Total training time: 0.05 seconds.
-- Epoch 4
Norm: 1.02, NNZs: 15, Bias: -0.658292, T: 150000, Avg. loss: 0.382083
Total training time: 0.06 seconds.
-- Epoch 5
Norm: 1.04, NNZs: 15, Bias: -0.719528, T: 187500, Avg. loss: 0.380486
Total training time: 0.07 seconds.
-- Epoch 6
Norm: 1.05, NNZs: 15, Bias: -0.763409, T: 225000, Avg. loss: 0.379578
Total training time: 0.09 seconds.
-- Epoch 7
Norm: 1.06, NNZs: 15, Bias: -0.795106, T: 262500, Avg. loss: 0.379150
Total training time: 0.10 seconds.
-- Epoch 8
Norm: 1.06, NNZs: 15, Bias: -0.819925, T: 300000, Avg. loss: 0.378856
Total training time: 0.11 seconds.
-- Epoch 9
Norm: 1.07, NNZs: 15, Bias: -0.837805, T: 337500, Avg. loss: 0.378585
Total training time: 0.13 seconds.
-- Epoch 10
Norm: 1.08, NNZs: 15, Bias: -0.853138, T: 375000, Avg. loss: 0.378630
Total training time: 0.14 seconds.
Convergence after 10 epochs took 0.14 seconds
```

```
Out[8]: SGDClassifier(alpha=0.0001, average=False, class_weight=None,
                    early_stopping=False, epsilon=0.1, eta0=0.0001,
                    fit_intercept=True, l1_ratio=0.15, learning_rate='constant',
                    loss='log', max_iter=1000, n_iter_no_change=5, n_jobs=None,
                    penalty='l2', power_t=0.5, random_state=15, shuffle=True,
                    tol=0.001, validation_fraction=0.1, verbose=2, warm_start=False)
```

```
In [9]: clf.coef_, clf.coef_.shape, clf.intercept_
#clf.coef_ will return the weights
#clf.coef_.shape will return the shape of weights
#clf.intercept_ will return the intercept term
```

```
Out[9]: (array([[ -0.42336692,  0.18547565, -0.14859036,  0.34144407, -0.2081867 ,
                  0.56016579, -0.45242483, -0.09408813,  0.2092732 ,  0.18084126,
                  0.19705191,  0.00421916, -0.0796037 ,  0.33852802,  0.02266721]]),
        (1, 15),
        array([ -0.8531383]))
```

This is formatted as code

Implement Logistic Regression with L2 regularization Using SGD: without using sklearn

1. We will be giving you some functions, please write code in that functions only.
2. After every function, we will be giving you expected output, please make sure that you get that output.

- Initialize the weight_vector and intercept term to zeros (Write your code in `def initialize_weights()`)
- Create a loss function (Write your code in `def logloss()`)

$$\text{logloss} = -1 * \frac{1}{n} \sum_{\text{foreach } Y_t, Y_{\text{pred}}} (Y_t \log_{10}(Y_{\text{pred}}) + (1 - Y_t) \log_{10}(1 - Y_{\text{pred}}))$$

- for each epoch:
 - for each batch of data points in train: (keep batch size=1)
 - calculate the gradient of loss function w.r.t each weight in weight vector (write your code in `def gradient_dw()`)

$$dw^{(t)} = x_n(y_n - \sigma((w^{(t)})^T x_n + b^t)) - \frac{\lambda}{N} w^{(t)}$$

- Calculate the gradient of the intercept (write your code in `def gradient_db()`) [check this](https://drive.google.com/file/d/1nQ08-XY4zvOLzRX-IGf8EYB5arb7-m1H/view?usp=sharing) (<https://drive.google.com/file/d/1nQ08-XY4zvOLzRX-IGf8EYB5arb7-m1H/view?usp=sharing>)

$$db^{(t)} = y_n - \sigma((w^{(t)})^T x_n + b^t)$$

- Update weights and intercept (check the equation number 32 in the above mentioned [pdf](https://drive.google.com/file/d/1nQ08-XY4zvOLzRX-IGf8EYB5arb7-m1H/view?usp=sharing) (<https://drive.google.com/file/d/1nQ08-XY4zvOLzRX-IGf8EYB5arb7-m1H/view?usp=sharing>)):
$$w^{(t+1)} \leftarrow w^{(t)} + \alpha(dw^{(t)})$$

$$b^{(t+1)} \leftarrow b^{(t)} + \alpha(db^{(t)})$$

- calculate the log loss for train and test with the updated weights (you can check the python assignment 10th question)
- And if you wish, you can compare the previous loss and the current loss, if it is not updating, then you can stop the training
- append this loss in the list (this will be used to see how loss is changing for each epoch after the training is over)

Initialize weights

```
In [10]: import math
```

```
In [11]: def initialize_weights(dim):  
    ''' In this function, we will initialize our weights and bias'''  
    #initialize the weights to zeros array of (1,dim) dimensions  
    #you use zeros_like function to initialize zero, check this link https://docs.scipy.org/doc/numpy/reference/generated/numpy.zeros\_like.html  
    #initialize bias to zero  
    w = np.zeros_like(dim)  
    b=0  
  
    return w,b
```

```
In [12]: dim=X_train[0]  
w,b = initialize_weights(dim)  
print('w =',(w))  
print('b =',str(b))  
  
w = [0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]  
b = 0
```

Grader function - 1

```
In [13]: dim=X_train[0]  
w,b = initialize_weights(dim)  
def grader_weights(w,b):  
    assert((len(w)==len(dim)) and b==0 and np.sum(w)==0.0)  
    return True  
grader_weights(w,b)
```

Out[13]: True

Compute sigmoid

$$\text{sigmoid}(z) = 1/(1 + \exp(-z))$$

```
In [14]: def sigmoid(z):  
    ''' In this function, we will return sigmoid of z'''  
    # compute sigmoid(z) and return  
    sigmoid = 1 / (1 + math.exp(-z))  
  
    return sigmoid
```

```
In [15]: sigmoid(2)
```

Out[15]: 0.8807970779778823

Grader function - 2

```
In [16]: def grader_sigmoid(z):
          val=sigmoid(z)
          assert(val==0.8807970779778823)
          return True
          grader_sigmoid(2)
```

Out[16]: True

Compute loss

$$\text{logloss} = -1 * \frac{1}{n} \sum_{\text{foreach } Y_t, Y_{\text{pred}}} (Y_t \log_{10}(Y_{\text{pred}}) + (1 - Y_t) \log_{10}(1 - Y_{\text{pred}}))$$

```
In [17]: def logloss(y_true,y_pred):
          '''In this function, we will compute log loss '''
          s = 0
          n = len(y_true)
          for i in range(0,n):
              s = s + ( (y_true[i]*math.log10(y_pred[i])) + (1-y_true[i])*math.log10(
1-y_pred[i]) )
              loss = -1 * (1/n)*s

          return loss
```

Grader function - 3

```
In [18]: def grader_logloss(true,pred):
          loss=logloss(true,pred)
          assert(loss==0.07644900402910389)
          return True
          true=[1,1,0,1,0]
          pred=[0.9,0.8,0.1,0.8,0.2]
          grader_logloss(true,pred)
```

Out[18]: True

Compute gradient w.r.to 'w'

$$dw^{(t)} = x_n(y_n - \sigma((w^{(t)})^T x_n + b^t)) - \frac{\lambda}{N} w^{(t)}$$

```
In [19]: def gradient_dw(x,y,w,b,alpha,N):
          '''In this function, we will compute the gardient w.r.to w '''

          dw =x*(y-sigmoid(np.dot(w.T,x)+b)) - ((alpha*w.T)/N)
          return dw
```

```
In [20]: print(w.T)
```

```
[0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
```

Grader function - 4

```
In [21]: def grader_dw(x,y,w,b,alpha,N):
          grad_dw=gradient_dw(x,y,w,b,alpha,N)
          assert(np.sum(grad_dw)==2.613689585)
          return True
          grad_x=np.array([-2.07864835,  3.31604252, -0.79104357, -3.87045546, -1.14783
          286,
                           -2.81434437, -0.86771071, -0.04073287,  0.84827878,  1.99451725,
                           3.67152472,  0.01451875,  2.01062888,  0.07373904, -5.54586092])
          grad_y=0
          grad_w,grad_b=initialize_weights(grad_x)
          alpha=0.0001
          N=len(X_train)
          grader_dw(grad_x,grad_y,grad_w,grad_b,alpha,N)
```

Out[21]: True

Compute gradient w.r.to 'b'

$$db^{(t)} = y_n - \sigma((w^{(t)})^T x_n + b^t)$$

```
In [22]: def gradient_db(x,y,w,b):
          '''In this function, we will compute gradient w.r.to b '''
          db = y - sigmoid(np.dot(w.T,x)+b)
          return db
```

Grader function - 5

```
In [23]: def grader_db(x,y,w,b):
          grad_db=gradient_db(x,y,w,b)
          assert(grad_db==-0.5)
          return True
          grad_x=np.array([-2.07864835,  3.31604252, -0.79104357, -3.87045546, -1.14783
          286,
                           -2.81434437, -0.86771071, -0.04073287,  0.84827878,  1.99451725,
                           3.67152472,  0.01451875,  2.01062888,  0.07373904, -5.54586092])
          grad_y=0
          grad_w,grad_b=initialize_weights(grad_x)
          alpha=0.0001
          N=len(X_train)
          grader_db(grad_x,grad_y,grad_w,grad_b)
```

Out[23]: True

Implementing logistic regression

In [26]:

w

```
Out[26]: array([-4.29755932e-01,  1.93023807e-01, -1.48464457e-01,  3.38103415e-01,
                -2.21228940e-01,  5.69932597e-01, -4.45183638e-01, -8.99209785e-02,
                 2.21804834e-01,  1.73809448e-01,  1.98727704e-01, -5.59450918e-04,
                -8.13106259e-02,  3.39094296e-01,  2.29784893e-02])
```

In [27]:

b

```
Out[27]: -0.8918925964701089
```

Goal of assignment

Compare your implementation and SGDClassifier's the weights and intercept, make sure they are as close as possible i.e difference should be in terms of 10^{-3}

```
In [28]: # these are the results we got after we implemented sgd and found the optimal weights and intercept
w-clf.coef_, b-clf.intercept_
```

```
Out[28]: (array([[ -0.00638902,  0.00754815,  0.0001259 , -0.00334065, -0.01304224,
                  0.00976681,  0.00724119,  0.00416715,  0.01253163, -0.00703181,
                  0.0016758 , -0.00477861, -0.00170693,  0.00056628,  0.00031128]]),
          array([-0.0387543]))
```

Plot epoch number vs train , test loss

- epoch number on X-axis
- loss on Y-axis

In [29]: tr_loss


```
Out[29]: [0.17546926223702297,  
0.16868174436540095,  
0.16639953379688238,  
0.16537404901928135,  
0.16486122004082515,  
0.16459114506307687,  
0.16444479874475598,  
0.16436411522525887,  
0.1643191231082826,  
0.1642938291559779,  
0.16427952013540825,  
0.16427138331585087,  
0.16426673469647732,  
0.1642640668101489,  
0.16426252835733005,  
0.16426163646238584,  
0.1642611161883861,  
0.16426081044856242,  
0.16426062918210865,  
0.16426052056735863,  
0.16426045466349845,  
0.16426041408844066,  
0.1642603886924852,  
0.1642603725071067,  
0.16426036199220195,  
0.1642603550261018,  
0.16426035032139918,  
0.16426034708556284,  
0.16426034482265406,  
0.16426034321669902,  
0.16426034206250442,  
0.16426034122417682,  
0.16426034060998718,  
0.16426034015685814,  
0.16426033982070054,  
0.16426033957023392,  
0.1642603393829793,  
0.16426033924262076,  
0.16426033913719898,  
0.16426033905789483,  
0.16426033899817297,  
0.1642603389531533,  
0.16426033891919317,  
0.1642603388935659,  
0.16426033887421562,  
0.16426033885960006,  
0.16426033884856042,  
0.16426033884022173,  
0.16426033883392135,  
0.16426033882915692]
```

In [30]: te_loss

```
Out[30]: [0.1759668786191598,  
0.16940989611779378,  
0.16721415304424353,  
0.16622329469756567,  
0.16572403546384085,  
0.16545876819806699,  
0.16531365222077019,  
0.16523283564551472,  
0.1651872786451111,  
0.16516136116063063,  
0.165146502653406,  
0.16513792321047355,  
0.16513293346948432,  
0.1651300087564692,  
0.16512827929561427,  
0.16512724616938745,  
0.16512662167682646,  
0.16512623900834086,  
0.16512600086609153,  
0.1651258501056387,  
0.1651257528922146,  
0.16512568899882157,  
0.16512564619492068,  
0.16512561698614076,  
0.16512559670985402,  
0.1651255824156623,  
0.16512557220222615,  
0.16512556482072047,  
0.16512555943510132,  
0.16512555547524393,  
0.165125552545619,  
0.16512555036754414,  
0.1651255487419863,  
0.16512554752515896,  
0.16512554661218476,  
0.1651255459259709,  
0.16512554540949317,  
0.16512554502035995,  
0.16512554472694344,  
0.16512554450556408,  
0.1651255443384607,  
0.16512554421227985,  
0.16512554411697833,  
0.16512554404498142,  
0.16512554399058435,  
0.16512554394947904,  
0.1651255439184143,  
0.16512554389493753,  
0.1651255438771939,  
0.16512554386378192]
```

```
In [33]: import matplotlib.pyplot as plt
ep=np.arange(1,epochs+1)
plt.plot(ep,tr_loss)
plt.plot(ep,te_loss)
plt.legend(['Train loss','Test loss'])
plt.xlabel('Epochs')
plt.ylabel('Loss Value')
plt.title('Epochs vs Loss Value')
```

Out[33]: Text(0.5, 1.0, 'Epochs vs Loss Value')



```
In [32]: def pred(w,b, X):
    N = len(X)
    predict = []
    for i in range(N):
        z=np.dot(w,X[i])+b
        if sigmoid(z) >= 0.5: # sigmoid(w,x,b) returns 1/(1+exp(-(dot(x,w)+
b)))
            predict.append(1)
        else:
            predict.append(0)
    return np.array(predict)
print(1-np.sum(y_train - pred(w,b,X_train))/len(X_train))
print(1-np.sum(y_test - pred(w,b,X_test))/len(X_test))
```

0.95224

0.95

In []:

In []: